

[DT0834]

Fundamentals and Applications of Artificial Intelligence

(Fondamenti e Applicazioni dell'Intelligenza Artificiale)

[F3I] B.Sc. Computer Science (L31 - *Laurea in Informatica*)

[2024/25] [Fall semester] [Teams: [9c7i3vd](#)]

2.1 – Intelligent agents

Prof. Francesco Gullo

University of L'Aquila, DISIM Department

francesco.gullo@univaq.it

<https://fgullo.github.io/>



2.1 Intelligent agents

~~1. Introduction to AI (6h)~~

- ~~- main notions, history, examples, applications~~

2. Agent-centric AI and symbolic AI (20h)

- intelligent agents, problem solving, solving problems by searching, constraint satisfaction problems, adversarial/multi-agent search, knowledge representation, automated reasoning, logic programming, examples and applications (game playing, solving NP-hard problems, robotics, autonomous cars, natural language processing, automated theorem proving, planning, expert systems)

3. Machine learning and connectionist AI (14h)

- supervised learning, regression and classification, neural networks and deep learning, examples and applications (generative AI, large language models, computer vision), other supervised-learning techniques (hints), other learning paradigms: unsupervised learning, semi-supervised learning, reinforcement learning (hints)

4. Neuro-symbolic AI (4h)

5. Philosophy, ethics and safety of AI (4h)



2.1 Intelligent agents - Outline

- Intelligent agents and agent-based AI tasks
- Rational agents
- Agent's Performance, Environment, Actuators, Sensors (PEAS)
- Agent function and agent program
- Main types of agent
- Example: paper buying agent

2.1 Intelligent agents

Materials for this lesson

Mostly:

- [BOOK] Russell, Norvig, 2021 - Artificial Intelligence: A Modern Approach, 4th ed. (<https://aima.cs.berkeley.edu/global-index.html>), Chapter 2

Further materials:

- [BOOK] Poole, Mackworth, 2023 - Artificial Intelligence: Foundations of Computational Agents, 3rd ed. (<https://artint.info/3e/html/ArtInt3e.html>), Chapters 1-2
- [COURSE] Introduction to Artificial Intelligence @UCBerkeley (<https://inst.eecs.berkeley.edu/~cs188/su24/>), Lectures 1-2
- [BOOK] Aggarwal, 2021 - Artificial Intelligence: A Textbook (<http://www.charuaggarwal.net>), Section 1.4

2.1 Intelligent agents

Materials for this lesson

Code:

- Companion Python code repository of Russell and Norvig's book ("Artificial Intelligence: A Modern Approach")
 - <https://github.com/aimacode/aima-python>
- Companion Python code repository of Poole and Mackworth's book ("Artificial Intelligence: Foundations of Computational Agents")
 - Web: <https://artint.info/AIPython/>
 - Code document: <https://artint.info/AIPython/aipython/aipython.pdf>
 - Code base: <https://artint.info/AIPython/aipython.zip>

Intelligent agents and agent-based AI tasks

Intelligent agents

Earlier AI approaches to action tended to focus on **single, isolated** software systems that acted in a relatively inflexible way, automatically following pre-set rules.

Modern technologies and software applications have created a need for **artificial entities** that are **more autonomous, flexible, and adaptive**, and that **operate as social entities** by interacting with each other and with the humans or the environment.

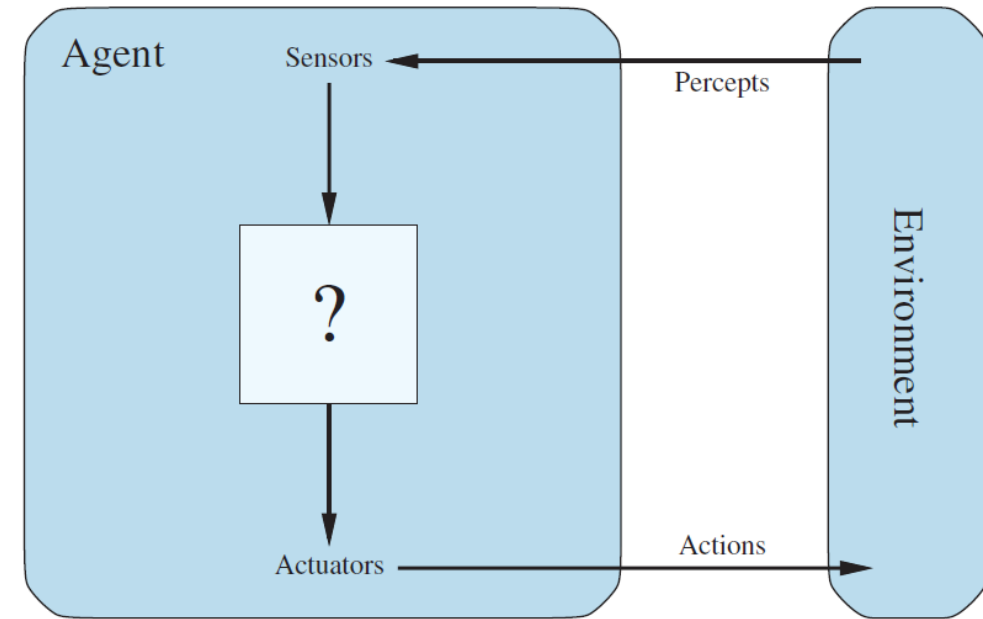
- Such artificial entities are termed “**intelligent agents**” or simply “**agents**”

Intelligent agents

The **concept of agent** is used **often** in AI, when a **sequence of decisions** (“**actions**”) need to be made in order to achieve a particular goal.

An **agent** is anything that can be viewed as **perceiving its environment through sensors** and **acting upon that environment through actuators**:

- A **human agent** has eyes, ears, and other organs for sensors and hands, legs, vocal tract, and so on for actuators.
- A **robotic agent** might have cameras and infrared range finders for sensors and various motors for actuators.
- A **software agent** receives file contents, network packets, and human input (keyboard/mouse/touchscreen/voice) as sensory inputs and acts on the environment by writing files, sending network packets, and displaying information or generating sounds.



[image source](#)

Intelligent agents: environment and perception

The **environment** which intelligent agents interact with could be everything—even the entire universe!

In practice it is just **that part of the universe** whose state we care about when designing this agent: **the part that affects what the agent perceives and that is affected by the agent's actions.**

The term “**percept**” is used to refer to the content an agent's sensors are perceiving.

An agent's **percept sequence** is the complete history of everything the agent has ever perceived.

In general, an **agent's choice of action** at any given instant can **depend on its built-in knowledge and on the entire percept sequence** observed to date, **but not on anything it hasn't perceived.**

Intelligent agents: examples

Tasks such as playing a game of chess, performing a robotic task, finding a path in a maze, replying to natural questions posed by a human, all require a sequence of decisions from an intelligent agent.

An example of an agent might be a robot, a chatbot, a chessplaying entity, or a route planner.

Intelligent agents: examples


Table 1.2: Agent examples

Agent	Sensor	Actuator	Objective	State	Environment
Cleaning Robot	Camera Joint sensor	Limbs Joints	Cleanliness evaluation	Object/joint positions	Cleaned space
Chess agent	Board input interface	Move output interface	Position evaluation	Chess position	Chess board
Self-driving car	Camera Sound sensor	Car control interface	Driving safety and goals	Speed/Position in traffic	Traffic conditions
Chatbot	Keyboard	Screen	Chat evaluation	Dialog history	Chat participants


[image source](#)

Intelligent agents: examples

Pac-Man is a registered trademark of Namco-Bandai Games, used here for educational purposes



Pac-Man is an agent!



[image source](#)

Agent-based AI tasks

It is only in the case of problems associated with a **sequence of decisions** that it becomes **important to use the concept of agent** and a state of the system.

Such settings are referred to as **sequential**.

An intelligent agent plays the same role that a human plays in a biological system,
based on a sequence of choices that achieve a desired goal.

Non-agent-based AI tasks

One does **not always encounter the notion of an agent in all forms of AI.**

For example, **the concept of agent rarely arises for single-decision problems** in machine learning; an example is that of **classification**, where one is trying to **categorize an instance into one of many classes** (e.g., categorize whether an image contains animals or not).

Such environmental settings are referred to as **episodic**, and even though **the notion of agent is implicit**, it is **rarely used in practice**.

In these cases, **the solution to this problem is a single-step process**, and therefore the question of system state or environment is not as important as it would be in a sequential process where an action affects subsequent states.

Therefore, in such cases, **the implicit agent receives a single percept and performs a single action**. The next percept-action pair is independent of the current one. In other words, **the individual episodes are independent of one another**.

Rational agents

Rationality and performance measure

A **rational agent** is one that “**does the right thing**”.

In the AI landscape, “**doing the right thing**” for an agent is typically recognized as corresponding to **how good the consequences yielded by the agent are**.

- This is called **consequentialism**.

More precisely, a **sequence of actions** performed by an agent causes the environment to go **through a sequence of states**. If such a **sequence of states is desirable**, then the agent has performed well (it has “**done the right thing**”).

This notion of **desirability** is captured by an ad-hoc defined **performance measure** that evaluates any given sequence of environment states.

Rationality vs. humanity

Humans have desires and preferences of their own, so **the notion of rationality as applied to humans** has to do with their success in choosing actions that produce sequences of environment states that are **desirable from their point of view**.

Machines, on the other hand, **do not have desires and preferences of their own**; **the performance measure** is, initially at least, in the **mind of the designer** of the machine, or in the **mind of the users** the machine is designed for.

Definition of a rational agent


*For each possible **percept sequence**, a **rational agent** should **select an action** that is **expected to maximize its performance measure**, **given** the evidence provided by **the percept sequence and whatever built-in knowledge** the agent has.*

Designing performance measures

It can be quite **hard** to formulate a **performance measure correctly**.

For instance, consider a simple **robotic vacuum-cleaning agent** in a world consisting of squares that can be either dirty or clean.

- Measure performance by the **amount of dirt** cleaned up in a single eight-hour shift => a rational agent can **maximize this performance measure** by cleaning up the dirt, then dumping it all on the floor, then cleaning it up again, and so on => **not exactly a good behavior!**
- A **more suitable performance measure** would reward the agent for having a **clean floor**. For example, one point could be awarded for each clean square at each time step (perhaps with a penalty for electricity consumed and noise generated).



[image source](#)


As a general rule, it is **better** to design performance measures according to **what one actually wants to be achieved in the environment**, rather than according to how one thinks the agent should behave.

Example of rational agent

Consider again the simple **vacuum-cleaner agent** that cleans a square if it is dirty and moves to the other square if not.

Assume the following:

- The **performance measure** awards one point for each clean square at each time step, over a “lifetime” of 1000 time steps, minus a **penalty** of <1 point for every action (for electricity consumed and noise generated).
- The available **actions** are **Right, Left, Suck, Do nothing**.
- The “**geography**” of the **environment** is **known a priori** but the dirt distribution and the initial location of the agent are not. Clean squares stay clean and sucking cleans the current square. The Right and Left actions move the agent one square except when this would take the agent outside the environment, in which case the agent remains where it is.
- The agent **correctly perceives its location** and **whether that location contains dirt**.



[image source](#)

A vacuum-cleaner agent that operates according to the **above rules** and tries to **maximize the above performance measure** is a **rational agent**.

Learning

It is **desirable** that a rational agent not only gather information but also to **learn** as much as possible **from what it perceives**.

The agent's initial configuration could reflect some prior knowledge of the environment, but **as the agent gains experience** this may be **modified and augmented**.

There are cases in which the **environment is completely known** a priori and completely predictable. **In such cases, the agent need not perceive or learn**; it simply acts correctly. But these are **extreme** and rather unlikely **cases**.

Designing an agent: Performance, Environment, Actuators, Sensors (PEAS)

PEAS: Performance, Environment, Actuators, Sensors

To **design an agent** we need to specify a “**PEAS**”, that is:

- **performance measure,**
- **environment,**
- **actuators,**
- **sensors.**

PEAS for automated taxi

Agent Type	Performance Measure	Environment	Actuators	Sensors
Taxi driver	Safe, fast, legal, comfortable trip, maximize profits, minimize impact on other road users	Roads, other traffic, police, pedestrians, customers, weather	Steering, accelerator, brake, signal, horn, display, speech	Cameras, radar, speedometer, GPS, engine sensors, accelerometer, microphones, touchscreen

Figure 2.4 PEAS description of the task environment for an automated taxi driver.

[image source](#)

PEAS for other agents

Agent Type	Performance Measure	Environment	Actuators	Sensors
Medical diagnosis system	Healthy patient, reduced costs	Patient, hospital, staff	Display of questions, tests, diagnoses, treatments	Touchscreen/voice entry of symptoms and findings
Satellite image analysis system	Correct categorization of objects, terrain	Orbiting satellite, downlink, weather	Display of scene categorization	High-resolution digital camera
Part-picking robot	Percentage of parts in correct bins	Conveyor belt with parts; bins	Jointed arm and hand	Camera, tactile and joint angle sensors

[image source](#)

PEAS for other agents

Refinery controller	Purity, yield, safety	Refinery, raw materials, operators	Valves, pumps, heaters, stirrers, displays	Temperature, pressure, flow, chemical sensors
Interactive English tutor	Student's score on test	Set of students, testing agency	Display of exercises, feedback, speech	Keyboard entry, voice

Figure 2.5 Examples of agent types and their PEAS descriptions.

[image source](#)

Types of environment

- Fully observable vs. partially observable
- Single agent vs. multi-agent
- Deterministic vs. nondeterministic
- Episodic vs. sequential
- Static vs. dynamic
- Discrete vs. continuous
- Known vs. unknown

Types of environment

Fully observable vs. partially observable

- Fully observable: if agent's sensors detect all aspects that are relevant to the choice of action; relevance, in turn, depends on the performance measure.
- An environment might be partially observable because of noisy or inaccurate sensors or because parts of the state are simply missing from the sensor data
 - For example, a vacuum agent with only a local dirt sensor cannot tell whether there is dirt in other squares, and an automated taxi cannot see what other drivers are thinking.

Single agent vs. multi-agent

- An agent solving a crossword puzzle by itself is clearly in a single-agent environment, whereas an agent playing chess is in a two-agent environment.

Types of environment

Deterministic vs. nondeterministic

- **Deterministic**: if the next state of the environment is **completely determined** by the current **state** and the **action** executed by the agent(s).
- Environment in **taxi driving is nondeterministic**, because one can never predict the behavior of traffic exactly; moreover, one's tires may blow out unexpectedly and one's engine may seize up without warning.
- The **vacuum-cleaning** world as we described **is deterministic**, but variations can include nondeterministic elements such as randomly appearing dirt and an unreliable suction mechanism.

“**Stochastic**” is sometimes used as a synonym for “**nondeterministic**”.

Here we make a **distinction** between these two terms:

- **Stochastic**: **probabilities** explicitly present (e.g., “there's a 25% chance of rain tomorrow”)
- **Nondeterministic**: **possibilities** are listed **without being quantified** (e.g., “there's a chance of rain tomorrow”).

Types of environment

Episodic vs. sequential

- **Episodic**: agent's experience is divided into **atomic episodes**. In each episode the agent receives a **percept** and then **performs a single action**. Crucially, **the next episode does not depend on the actions taken in previous episodes**.
 - **Many classification tasks are episodic**. For example, an agent that has to spot defective parts on an assembly line bases each decision on the current part, regardless of previous decisions; moreover, the current decision doesn't affect whether the next part is defective.
- **Sequential**: the **current decision** could **affect all future decisions**.
 - **Chess and taxi driving are sequential**: in both cases, short-term actions can have long-term consequences.

Episodic environments are **much simpler** than sequential environments because **the agent does not need to think ahead**.

Types of environment

Discrete vs. continuous

- The discrete/continuous distinction applies to the **state of the environment**, to the **way how time is handled**, and to the **percepts and actions of the agent**.
- For example, the **chess** environment has a **finite number of distinct states** (excluding the clock).
- **Chess** also has a **discrete set of percepts and actions**.
- **Taxi driving is a continuous-state and continuous-time problem**: the speed and location of the taxi and of the other vehicles sweep through a range of continuous values and do so smoothly over time.
- **Taxi-driving actions are also continuous** (steering angles, etc.). **Input from digital cameras is discrete**, strictly speaking, but is typically **treated as representing continuously varying intensities and locations**.

Types of environment

Known vs. unknown

- Strictly speaking, this distinction refers not to the environment itself but to the **agent's** (or designer's) **state of knowledge** about the “**laws of physics**” of the environment.
- In a **known** environment, the **outcomes** (or outcome probabilities if the environment is nondeterministic) **for all actions are given**.
- If the **environment is unknown**, the **agent** will have to **learn how it works** in order to make good decisions.

The distinction **between known and unknown** environments is **not the same as the one between fully and partially observable** environments.

- It is quite possible for a **known environment to be partially observable**—for example, in **solitaire card games**, I know the rules but am still unable to see the cards that have not yet been turned over.
- Conversely, an **unknown environment can be fully observable**—in a **new video game**, the screen may show the entire game state but I still don't know what the buttons do until I try them

Types of environment

Task Environment	Observable	Agents	Deterministic	Episodic	Static	Discrete
Crossword puzzle	Fully	Single	Deterministic	Sequential	Static	Discrete
Chess with a clock	Fully	Multi	Deterministic	Sequential	Semi	Discrete
Poker	Partially	Multi	Stochastic	Sequential	Static	Discrete
Backgammon	Fully	Multi	Stochastic	Sequential	Static	Discrete
Taxi driving	Partially	Multi	Stochastic	Sequential	Dynamic	Continuous
Medical diagnosis	Partially	Single	Stochastic	Sequential	Dynamic	Continuous
Image analysis	Fully	Single	Deterministic	Episodic	Semi	Continuous
Part-picking robot	Partially	Single	Stochastic	Episodic	Dynamic	Continuous
Refinery controller	Partially	Single	Stochastic	Sequential	Dynamic	Continuous
English tutor	Partially	Multi	Stochastic	Sequential	Dynamic	Discrete

[image source](#)

Figure 2.6 Examples of environments and their characteristics.

Perhaps the **hardest** case is *partially observable*, *multiagent*, *nondeterministic*, *sequential*, *dynamic*, *continuous*, and *unknown*.

Designing an agent: agent program

Agent function and agent program

Formally speaking, the behavior of an agent behavior is described by an **agent function** that **maps any given percept sequence to an action**.

The **agent function** is an **abstract mathematical description** of agent's behavior.

Internally to an agent, the **agent function is implemented** by what is called **agent program**.

- The **agent program** is a **concrete implementation of the agent function**, running within some physical system.

Agent architecture

The agent program is assumed to run on some sort of **computing device** with **physical sensors and actuators**: this is called “the **agent architecture**”

The architecture might be just an **ordinary PC**, or it might be a **robotic car** with several onboard computers, cameras, and other sensors.

In general, the architecture makes the **percepts from the sensors available to the program**, **runs the program**, and **feeds program's action choices to the actuators** as they are generated.

The **job of AI** is mainly to **design agent programs**, while **agent architectures** are somehow given for granted.

Agent program and agent architecture

The job of AI is to design an **agent program** that **implements the agent function**, i.e., the mapping from percepts to actions.

The agent program is assumed to run on some sort of **computing device** with **physical sensors and actuators**: this is called “the **agent architecture**”

Simply speaking, an agent is a combination of its program and its architecture:

$$\textit{agent} = \textit{program} + \textit{architecture}$$

Table-driven agent

Perhaps the **simplest** way to design an agent program is to build a **lookup table** that maps **every possible percept sequence to** the corresponding **most desirable action** to be performed.

```
function TABLE-DRIVEN-AGENT(percept) returns an action
  persistent: percepts, a sequence, initially empty
               table, a table of actions, indexed by percept sequences, initially fully specified


  append percept to the end of percepts
  action ← LOOKUP(percepts, table)
  return action
```

Figure 2.7 The TABLE-DRIVEN-AGENT program is invoked for each new percept and returns an action each time. It retains the complete percept sequence in memory.

[image source](#)

Table-driven agent

A simple table-driven vacuum-cleaner agent:



A vacuum-cleaner environment with just two locations. Each location can be clean or dirty, and the agent can move left or right and can clean the square that it occupies.

Percept sequence	Action
$[A, \textit{Clean}]$	<i>Right</i>
$[A, \textit{Dirty}]$	<i>Suck</i>
$[B, \textit{Clean}]$	<i>Left</i>
$[B, \textit{Dirty}]$	<i>Suck</i>
$[A, \textit{Clean}], [A, \textit{Clean}]$	<i>Right</i>
$[A, \textit{Clean}], [A, \textit{Dirty}]$	<i>Suck</i>
\vdots	\vdots
$[A, \textit{Clean}], [A, \textit{Clean}], [A, \textit{Clean}]$	<i>Right</i>
$[A, \textit{Clean}], [A, \textit{Clean}], [A, \textit{Dirty}]$	<i>Suck</i>
\vdots	\vdots

Partial tabulation of a simple agent function for the vacuum-cleaner world shown on the left. The agent cleans the current square if it is dirty, otherwise it moves to the other square. Note that the table is of **unbounded size** unless there is a restriction on the length of possible percept sequences.

Table-driven agent

Building a table-driven agent is **feasible** only in a **very few simple cases**.

In most cases, the lookup table is going to be of **huge** (or even **infinite**) **size**.

Let P be the set of **possible percepts** and let T be the **lifetime of the agent** (the total number of percepts it will receive).

The lookup table will contain $\sum_{t=1}^T |P|^t$ **entries**.

The **lookup table of chess** has at least 10^{150} **entries**.

- Much **more than the total number of atoms in the observable universe**, which is estimated to be $<10^{82}$!

Going beyond table-driven agents

The **key challenge for AI** is to find out how to write agent programs that, to the maximum extent possible, produce **rational behavior from a smallish program** rather than from a vast table.

For instance, a **non-table-driven vacuum-cleaner agent program** might be:

```
function REFLEX-VACUUM-AGENT([location,status]) returns an action
    if status = Dirty then return Suck
    else if location = A then return Right
    else if location = B then return Left
```

[image source](#)

Types of agent

Simple reflex agent

- takes actions based on the current percept only

Model-based reflex agent

- maintains internal state to track aspects of the world that are not evident in the current percept; takes actions based on the current percept and their internal state

Goal-based agent

- acts to achieve their goals

Utility-based agent

- tries to maximize their own expected “happiness”

All agents can **improve** their performance through **learning**.

Types of agent: simple reflex agent

Taxi-driving agent:

It brakes based on the **condition-action rule** “*if car-in-front-is-braking then initiate-braking*”.

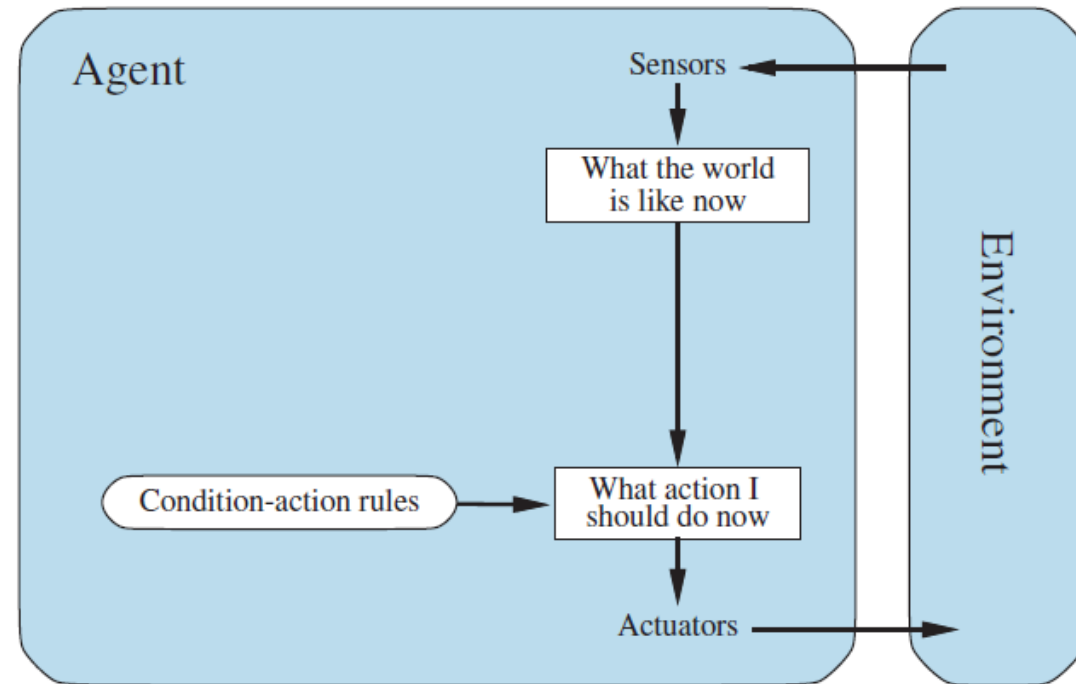


Figure 2.9 Schematic diagram of a simple reflex agent. We use rectangles to denote the current internal state of the agent’s decision process, and ovals to represent the background information used in the process.

[image source](#)

Types of agent: model-based reflex agent

Taxi-driving agent:

For driving tasks other than braking, such as **changing lanes**, the agent needs to **keep track of where the other cars are** (if it can't see them all at once).

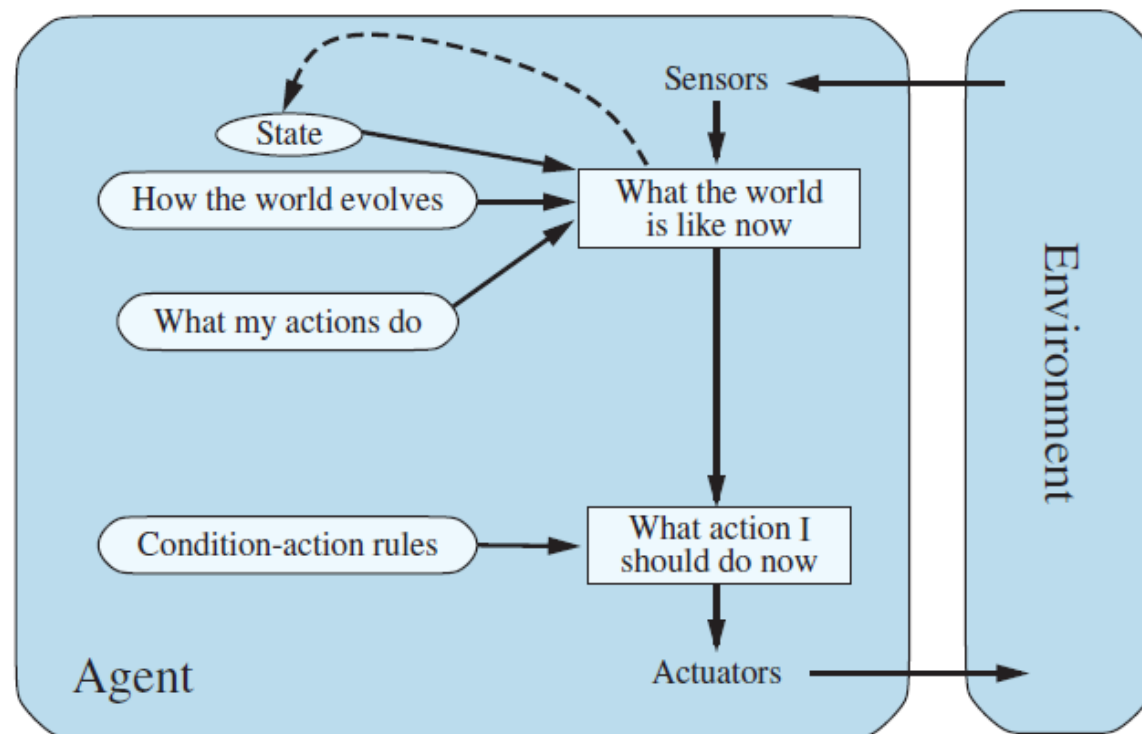


Figure 2.11 A model-based reflex agent.

[image source](#)

Types of agent: goal-based agent

Taxi-driving agent:

Knowing something about the current state of the environment is not always enough to decide what to do.

For example, at a road junction, the taxi can turn left, turn right, or go straight on.

The correct decision depends on its **goal**, i.e., on **where the taxi is trying to get to**.

Sometimes goal-based action selection is straightforward, i.e., when goal satisfaction results immediately from a single action.

Sometimes it will be more tricky, i.e., when the agent has to **consider long sequences** of twists and turns in order to find a way to achieve the goal.

Search and **planning** are the **subfields of AI** devoted to finding action sequences that achieve the agent's goals.

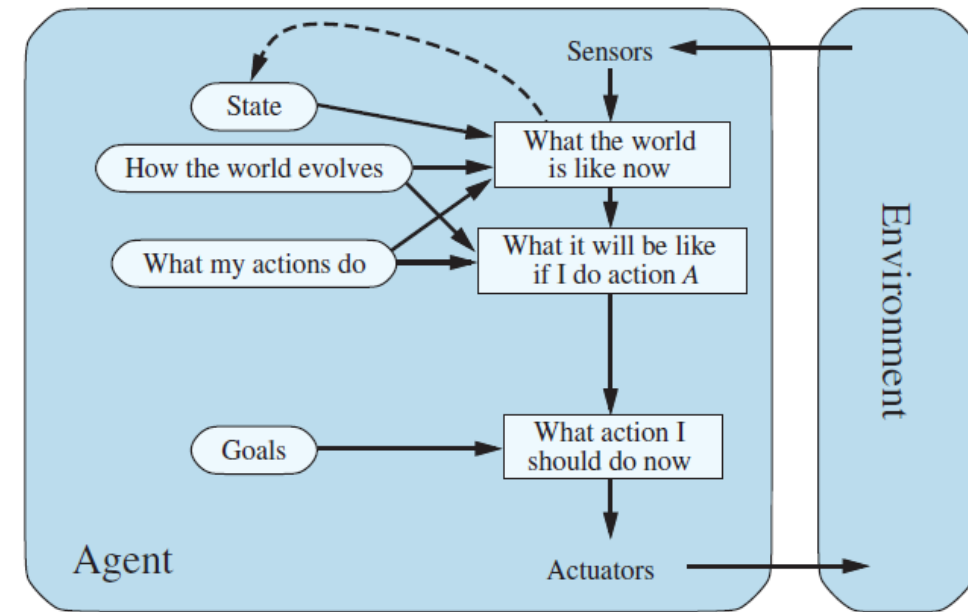


Figure 2.13 A model-based, goal-based agent. It keeps track of the world state as well as a set of goals it is trying to achieve, and chooses an action that will (eventually) lead to the achievement of its goals.

[image source](#)

Types of agent: utility-based agent

Taxi-driving agent:

Goals alone are not enough to generate high-quality behavior in most environments: many action sequences will get the taxi to its destination (thereby achieving the goal), but some are quicker, safer, more reliable, or cheaper than others.

Goals just provide a crude binary distinction between “happy” and “unhappy” states.

A more general **performance measure** should allow a comparison of different **environment states** according to exactly **how happy they would make the agent**.

The measure of “**happiness**” of an agent is mathematically expressed by a so-called **utility function**.

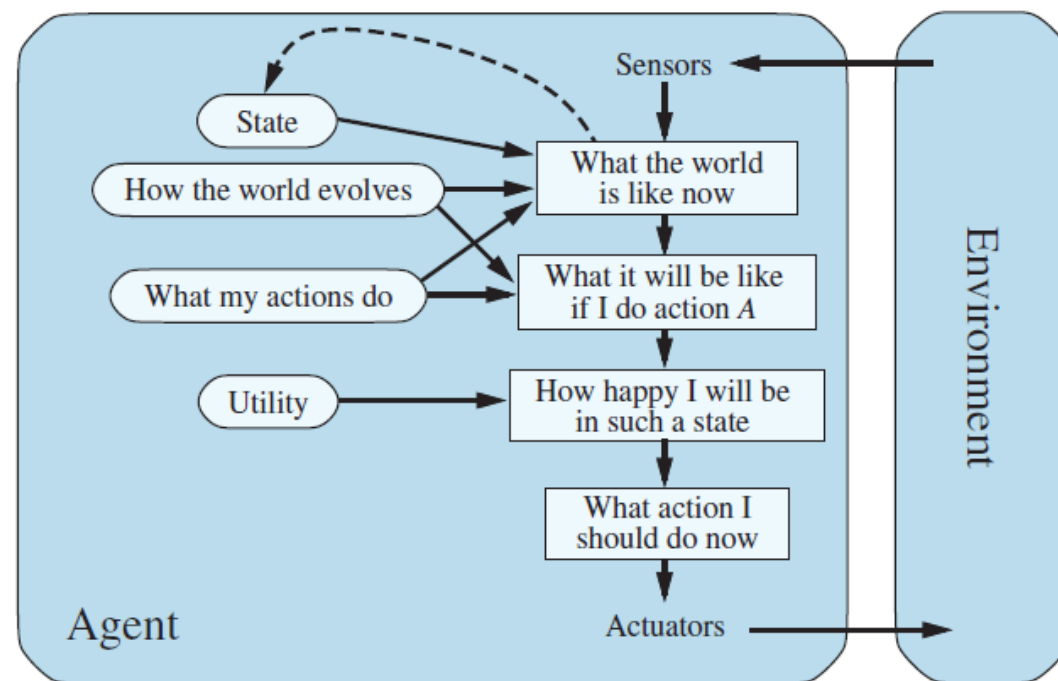


Figure 2.14 A model-based, utility-based agent. It uses a model of the world, along with a utility function that measures its preferences among states of the world. Then it chooses the action that leads to the best expected utility, where expected utility is computed by averaging over all possible outcome states, weighted by the probability of the outcome.

[image source](#)

Types of agent: utility-based agent

Utility function vs. performance measure:

- Agent's **utility function** is essentially an **internalization of the performance measure**.
- Provided that the **internal utility function** and the **external performance measure** are in **agreement**, an agent that chooses actions to **maximize its utility** will be **rational** according to the external performance measure.
 - Designing an agent like this **is not the only way** of achieving **rationality**!

Types of agent: learning agents

A **learning agent** in AI is the type of agent that can **learn from its past experiences**.

It **starts to act with basic knowledge** and then is able to act and **adapt automatically through learning**.

Any type of agent (model-based, goal-based, utility-based, etc.) can be built **as a learning agent (or not)**.

Historical note:

Turing, in **1950**, considered the idea of actually programming his intelligent machines by hand. He estimates how much work this might take and concludes, “*Some more expeditious method seems desirable*”.

The method he proposes is to **build learning machines and then to teach them**. In many areas of AI, this is now the preferred method for creating state-of-the-art systems.

Types of agent: learning agents

Taxi-driver agent:

Suppose it receives no tips from passengers who have been thoroughly shaken up during the trip.

The external performance standard must inform the agent that the loss of tips is a negative contribution to its overall performance.

Then the agent might be able to learn that violent maneuvers do not contribute to its own utility.

In a sense, the performance standard distinguishes part of the incoming percept as a **reward** (or **penalty**) that provides direct feedback on the quality of the agent's behavior.

This **reward/penalty** can be **used** by the agent **to learn** about good and bad behaviors.



Example: paper buying agent

Paper buying agent

Source:

- **Example 2.1** of Poole and Mackworth's book (*Artificial Intelligence: Foundations of Computational Agents*), 3rd ed.
- **utilities.py**, **display.py**, **agents.py**, and **agentBuying.py** files in companion Python code repository of Poole and Mackworth's book
 - <https://artint.info/AIPython/aipython.zip>
 - Section 2.2 of <https://artint.info/AIPython/aipython/aipython.pdf>

Paper buying agent

Agent responsible for guaranteeing that the available amount of some **household commodity** (e.g., **toilet paper**) is adequate.

It **monitors** (on a **daily basis**):

- **online deals** for buying paper;
- how much **paper** the household has **in stock**.


It **decides** (on a **daily basis**) whether to **order more paper** and **how much to order**.

Percepts:

- **price** of the paper, and
- **amount** of paper currently **in stock**.


Action: number of **units of paper** to be **ordered** (zero if the agent decides not to order any).

Paper buying agent

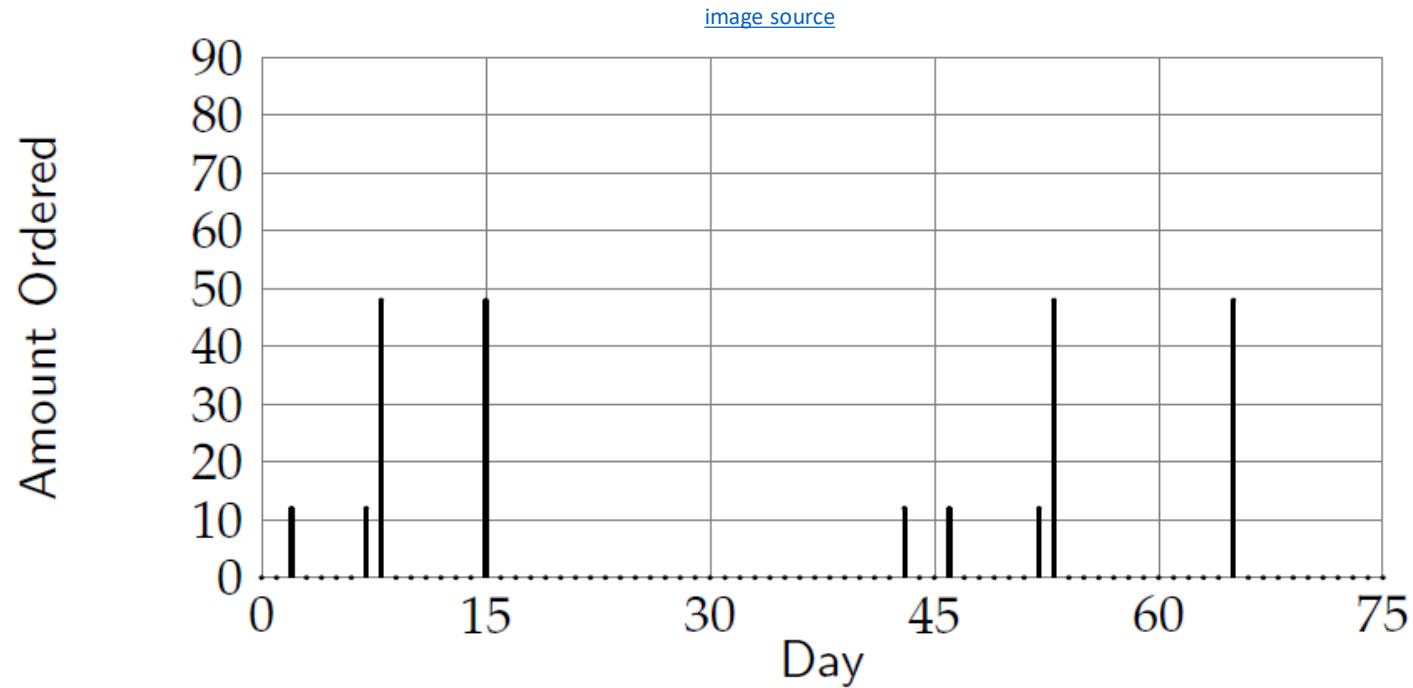


[image source](#)

A possible **percept sequence**
(for the two percepts,
namely price and amount in stock)



Paper buying agent



A possible **sequence of actions** performed

Paper buying agent: `display.py`

```
# display.py - A simple way to trace the intermediate steps of algorithms.

11 class Displayable(object):
12     """Class that uses 'display'.
13     The amount of detail is controlled by max_display_level
14     """
15     max_display_level = 1 # can be overridden in subclasses or instances
16
17     def display(self, level, *args, **nargs):
18         """print the arguments if level is less than or equal to the
19         current max_display_level.
20         level is an integer.
21         the other arguments are whatever arguments print can take.
22         """
23         if level <= self.max_display_level:
24             print(*args, **nargs) ##if error you are using Python2 not Python3
```

More on `*args` and `**nargs`: <https://www.linkedin.com/pulse/decoding-python-functions-default-positional-keyword-benjamin/>

Paper buying agent: `agents.py`

```
11 from display import Displayable
12
13 class Agent(Displayable):
14
15     def initial_action(self, percept):
16         """return the initial action."""
17         return self.select_action(percept) # same as select_action
18
19     def select_action(self, percept):
20         """return the next action (and update internal state) given percept
21         percept is variable:value dictionary
22         """
23         raise NotImplementedError("go") # abstract method
```

An agent takes the percept, updates its internal state (if any), and output its next action.

An agent implements the method **`select_action`** that takes percept and returns its next action.

Paper buying agent: `agents.py`

```
24
25 class Environment(Displayable):
26     def initial_percept(self):
27         """returns the initial percept for the agent"""
28         raise NotImplementedError("initial_percept") # abstract method
29
30     def do(self, action):
31         """does the action in the environment
32         returns the next percept """
33         raise NotImplementedError("Environment.do") # abstract method
```

An environment takes in actions of the agents, updates its internal state and returns the next percept, using the method `do`.

The environment implements a `do(action)` method where action is a variable-value dictionary. This returns a percept, which is also a variable-value dictionary. The use of dictionaries allows for structured actions and percepts.

Paper buying agent: `agents.py`

```
34
35 class Simulate(Displayable):
36     """simulate the interaction between the agent and the environment
37     for n time steps.
38     Returns a pair of the agent state and the environment state.
39     """
40     def __init__(self, agent, environment):
41         self.agent = agent
42         self.env = environment
43         self.percept = self.env.initial_percept()
44         self.percept_history = [self.percept]
45         self.action_history = []
46
47     def go(self, n):
48         for i in range(n):
49             action = self.agent.select_action(self.percept)
50             self.display(2, f"i={i} action={action}")
51             self.percept = self.env.do(action)
52             self.display(2, f"percept={self.percept}")
53
```

The methods `do` and `select_action` are chained together to build a simulator.

In order to start this, we need either an action or a percept.

There are two variants used:

- An agent implements the `initial_action` method which is used initially.
- The environment implements the `initial_percept` method which gives the initial percept. This is the method used here.

Paper buying agent: agentBuying.py

```
11 import random
12 from agents import Agent, Environment, Simulate
13 from utilities import select_from_dist
14
15 class TP_env(Environment):
16     price_delta = [0, 0, 0, 21, 0, 20, 0, -64, 0, 0, 23, 0, 0, 0, -35,
17                   0, 76, 0, -41, 0, 0, 0, 21, 0, 5, 0, 5, 0, 0, 0, 5, 0, -15, 0, 5,
18                   0, 5, 0, -115, 0, 115, 0, 5, 0, -15, 0, 5, 0, 5, 0, 0, 0, 5, 0,
19                   -59, 0, 44, 0, 5, 0, 5, 0, 0, 0, 5, 0, -65, 50, 0, 5, 0, 5, 0, 0,
20                   0, 5, 0]
21     sd = 5 # noise standard deviation
22
23     def __init__(self):
24         """paper buying agent"""
25         self.time=0
26         self.stock=20
27         self.stock_history = [] # memory of the stock history
28         self.price_history = [] # memory of the price history
```

The environment state is given in terms of the time and the amount of paper in stock.

It also remembers the in-stock history and the price history.

The percept consists of the price and the amount of paper in stock.

The action of the agent is the number of items to buy.

Paper buying agent: agentBuying.py

```
29
30     def initial_percept(self):
31         """return initial percept"""
32         self.stock_history.append(self.stock)
33         self.price = round(234+self.sd*random.gauss(0,1))
34         self.price_history.append(self.price)
35         return {'price': self.price,
36               'instock': self.stock}
37
38     def do(self, action):
39         """does action (buy) and returns percept consisting of price and instock"""
40         used = select_from_dist({6:0.1, 5:0.1, 4:0.1, 3:0.3, 2:0.2, 1:0.2})
41         # used = select_from_dist({7:0.1, 6:0.2, 5:0.2, 4:0.3, 3:0.1, 2:0.1}) # uses more paper
42         bought = action['buy']
43         self.stock = self.stock+bought-used
44         self.stock_history.append(self.stock)
45         self.time += 1
46         self.price = round(self.price
47                           + self.price_delta[self.time%len(self.price_delta)] # repeating pattern
48                           + self.sd*random.gauss(0,1)) # plus randomness
49         self.price_history.append(self.price)
50         return {'price': self.price,
51               'instock': self.stock}
```

Paper buying agent: agentBuying.py

```
52
53 class TP_agent(Agent):
54     def __init__(self):
55         self.spent = 0
56         percept = env.initial_percept()
57         self.ave = self.last_price = percept['price']
58         self.instock = percept['instock']
59         self.buy_history = []
60
```

Paper buying agent: agentBuying.py

```
61     def select_action(self, percept):
62         """return next action to carry out
63         """
64         self.last_price = percept['price']
65         self.ave = self.ave+(self.last_price-self.ave)*0.05
66         self.instock = percept['instock']
67         if self.last_price < 0.9*self.ave and self.instock < 60:
68             tobuy = 48
69         elif self.instock < 12:
70             tobuy = 12
71         else:
72             tobuy = 0
73         self.spent += tobuy*self.last_price
74         self.buy_history.append(tobuy)
75         return {'buy': tobuy}
```

Paper buying agent: `agentBuying.py`

```
76
77  env = TP_env()
78  ag = TP_agent()
79  sim = Simulate(ag, env)
80  sim.go(90)
81  ag.spent/env.time    ## average spent per time period
```


Paper buying agent: `utilities.py`

```
1 # utilities.py - AIPython useful utilities
2 # AIFCA Python code Version 0.9.13 Documentation at https://aipython.org
3 # Download the zip file and read aipython.pdf for documentation
4
5 # Artificial Intelligence: Foundations of Computational Agents https://artint.info
6 # Copyright 2017-2024 David L. Poole and Alan K. Mackworth
7 # This work is licensed under a Creative Commons
8 # Attribution-NonCommercial-ShareAlike 4.0 International License.
9 # See: https://creativecommons.org/licenses/by-nc-sa/4.0/deed.en
```

```
49 def select_from_dist(item_prob_dist):
50     """ returns a value from a distribution.
51     item_prob_dist is an item:probability dictionary, where the
52     probabilities sum to 1.
53     returns an item chosen in proportion to its probability
54     """
55     ranreal = random.random()
56     for (it,prob) in item_prob_dist.items():
57         if ranreal < prob:
58             return it
59         else:
60             ranreal -= prob
61     raise RuntimeError(f"{item_prob_dist} is not a probability distribution")
```


Questions?